

AI Metrics Dashboard

AI attribution lives on the PR and rolls up through story → epic → deliverable. "Done" = merged; cycle time is dev-side only.

Reporting period
Q3 2026 · Sprints 21–26

AI adoption, current sprint ⓘ

80%

4 of 5 devs · up from 40% in sprint 21

PR throughput, current sprint ⓘ

14/wk

vs. pre-AI baseline: 9/wk

Avg AI-assisted % across deliverables ⓘ

58%

share of stories w/ an AI-assisted PR

Say-do ratio, current sprint ⓘ

88%

baseline (pre-AI, 6-sprint avg): 77%

Escaped defect rate ⓘ


6%


AI-assisted 4% · manual 9%

Deliverable speed

Story points ÷ 4 pts/day baseline (pre-AI, 5-dev team velocity) vs. actual days taken


Deliverable	Story points	Est. days	Actual days	Faster by	AI-assisted %
Checkout redesign	84	21	12	9d faster (43%)	71%
Onboarding flow	80	20	11	9d faster (45%)	66%
Mobile push notifs	24	6	3	3d faster (50%)	82%
Admin permissions v2	52	13	9	4d faster (31%)	55%
Search relevance tune	16	4	4	On pace	44%
Billing migration	60	15	17	2d slower (13%)	38%






PR ledger
Attribution source of truth — every other rollup traces back to this table

PR	Story	Opened	Merged	Review	Reverted	AI
#2481	ENG-4021	Sep 08	Sep 09	2	No	Assisted
#2477	ENG-4018	Sep 04	Sep 07	5	No	Assisted
#2465	ENG-4012	Sep 02	Sep 05	6	No	Manual
#2460	ENG-4009	Aug 30	Aug 31	1	No	Assisted
#2452	ENG-4005	Aug 26	Sep 01	8	Reverted	Manual
#2441	ENG-3998	Aug 22	Aug 23	2	No	Assisted
#2436	ENG-3991	Aug 19	Aug 21	4	No	Manual



Say-do ratio by sprint
Completed ÷ committed story points, high vs low AI-assist sprints



Sample data for illustration. AI adoption = distinct PR authors with ≥ 1 AI-assisted PR merged in the sprint $\div$ total developers on the team (4 of 5 = 80%). This requires an author field on each PR record — GitHub provides it on every PR, so no extra instrumentation is needed. It is a share of people, not of PRs: a team can show a high AI-assisted PR share driven by two heavy users while most of the team never adopts, and only the people-based figure distinguishes those cases. Baseline pace = 4 story points/day, from the team's pre-AI velocity (40 points per 5-dev, 10-day sprint). Est. days = story points $\div$ baseline pace; "faster by" compares that estimate to actual days taken. PR throughput = merged PRs per sprint, split by whether the PR carries the AI-assisted flag — the plain count of work items closed, with no normalization or modeling. Story cycle time = earliest linked PR open $\rightarrow$ latest linked PR merge. A story is "AI-assisted" if any linked PR carries the flag. Deliverable AI-assisted % = share of child stories flagged AI-assisted (distinct from the PR-level yes/no flag it rolls up from). "Merged" is the Done event used here — it does not include QA, staging, or production deploy; track merge $\rightarrow$ deploy separately if you need that number. Defect root cause is attributed to the specific PR that introduced the offending code (via git blame on the fix, mapped back to a PR), not to the story's overall AI flag.